

Base Model

In [1]:

```
%load_ext autoreload
%autoreload 2
```

In [26]:

```
import os
import sys
from shutil import copyfile

import numpy as np
from PIL import Image
import matplotlib.pyplot as plt
import matplotlib.patches as patches
import torch
import torch.optim as optim
from torch.utils.data import TensorDataset, DataLoader
from torch_snippets import Report
import time
```

In [5]:

```
# ../src/base-model/
module_path = os.path.abspath(os.path.join('..../src/base-model/'))
if module_path not in sys.path:
    sys.path.append(module_path)

import utils
import rcnn
```

In [6]:

```
## load annotations of 135 images
anno_dict = np.load('../data/anno-big.npy', allow_pickle='TRUE').item()
imgs_person = list(anno_dict.keys())
```

In [7]:

```
len(imgs_person)
```

Out[7]:

135

In [8]:

```
imgs_path = '/floyd/home/datasets/base/'
```

In [9]:

```
# ## move the images to separate dir
# cityscapes_path = '/home/marko/data/cityscapes/leftImg8bit_trainvaltest/leftImg8bit/train/'
# imgs_path = '/home/marko/data/base/'

# for img_name in imgs_person:
#     city_name = img_name.split('_')[0]
#     img_src_path = cityscapes_path + city_name + '/' + img_name
#     copyfile(img_src_path, imgs_path + img_name)
```

In [10]:

```
i = 0
img_name = imgs_person[i]
img_path = imgs_path + img_name
img = Image.open(img_path)
bboxes = anno_dict[img_name]
```

In [11]:

```
plt.rcParams['figure.figsize'] = [12, 8]
```

In [12]:

```
fig, ax = plt.subplots()
ax.imshow(img)

for bbox in bboxes:
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='r', facecolor='none')

    ax.add_patch(rect)

plt.title(img_name)
plt.show()
```



In [13]:

```
regions = utils.get_simple_regions()
```

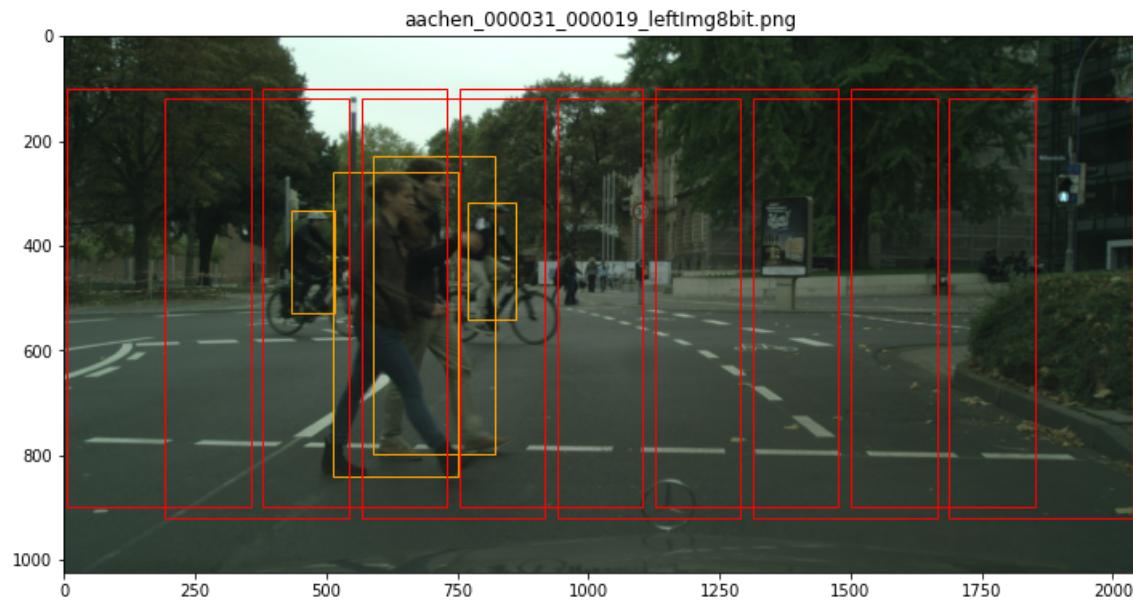
In [14]:

```
fig, ax = plt.subplots()
ax.imshow(img)

for bbox in bboxes:
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='orange', facecolor='none')
    ax.add_patch(rect)

for bbox in regions:
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='red', facecolor='none')
    ax.add_patch(rect)

plt.title(img_name)
plt.show()
```



In [15]:

```
img_np = np.array(img)
H, W = img_np.shape[:2]
```

In [16]:

```
H,W
```

Out[16]:

```
(1024, 2048)
```

Prepare data

In [17]:

```
all_img_names, all_labels, all_diffs, all_rois, all_bboxes = utils.get_data(
    imgs_person, anno_dict, regions, W, H)
```

In [18]:

```
len(all_img_names), len(all_labels)
```

Out[18]:

```
(135, 135)
```

In [19]:

```
## check results
i = 0
img_name = all_img_names[i]
bboxes = all_bboxes[i]

img_path = imgs_path + img_name
img = Image.open(img_path)

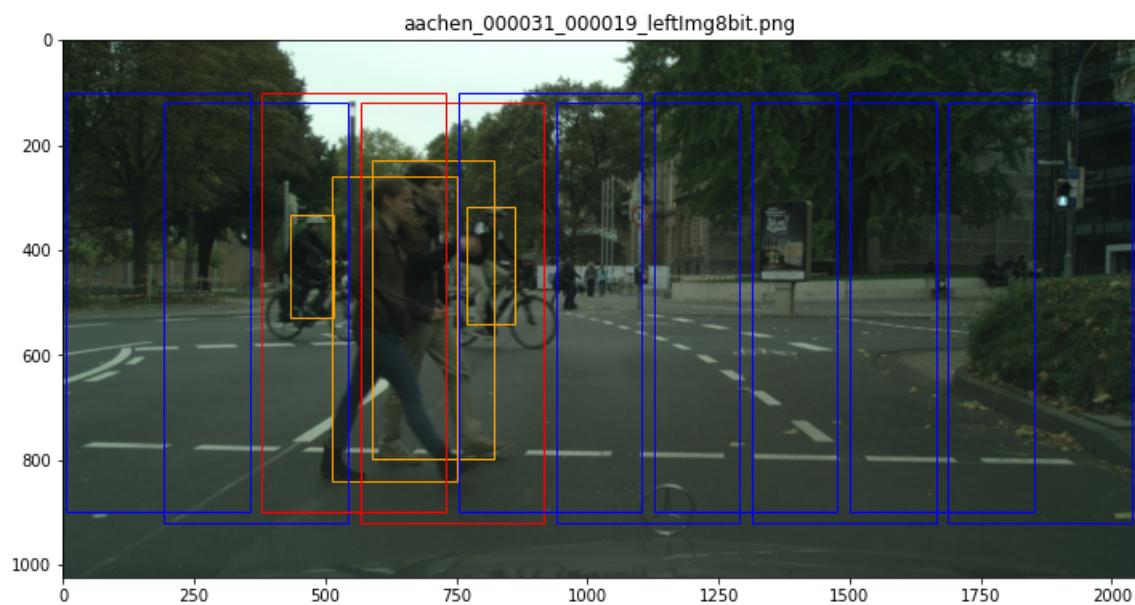
fig, ax = plt.subplots()
ax.imshow(img)

labels = all_labels[i]
colors = ['b', 'r']

for i, bbox in enumerate(bboxes):
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='orange', facecolor='none')
    ax.add_patch(rect)

for i, bbox in enumerate(regions):
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor=colors[labels[i]], facecolor='none')
    ax.add_patch(rect)

plt.title(img_name)
plt.show()
```



In [20]:

```
img_region = np.array([W, H, W, H]) # for scaling
diffs0 = all_diffs[0] * img_region
diffs0
```

Out[20]:

```
array([[-583., -128., 117., 231.],
       [-242., -212., 270., 604.],
       [-134., -158., 111., 216.],
       [-22., -108., 117., 231.],
       [165., -128., 117., 231.],
       [352., -108., 117., 231.],
       [539., -128., 117., 231.],
       [726., -108., 117., 231.],
       [913., -128., 117., 231.],
       [1100., -108., 117., 231.]])
```

In [21]:

```
i = 0
img_name = all_img_names[i]
bboxes = all_bboxes[i]
img_path = imgs_path + img_name
img = Image.open(img_path)

fig, ax = plt.subplots()
ax.imshow(img);

labels = all_labels[i]
colors = ['b', 'r']

# bbox = [x, y, w, h]
for i, bbox in enumerate(bboxes):
    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor='orange', facecolor='none')
    ax.add_patch(rect)

for i, bbox in enumerate(regions):
    if labels[i] == 1:
        bbox = bbox - diffs0[i] # add the diffs

    rect = patches.Rectangle(
        (bbox[0], bbox[1]), bbox[2], bbox[3],
        linewidth=1, edgecolor=colors[labels[i]], facecolor='none')
    ax.add_patch(rect)

plt.title(img_name)
plt.show()
```



Prepare train and test sets

In [22]:

```
img_paths = [(img_path + img_name) for img_name in all_img_names]
n_train = 8 * len(img_paths) // 10
n_train
```

Out[22]:

108

In [23]:

```
## img_paths, rois, labels, diffs, bboxes
train_set = rcnn.Dataset(
    img_paths[:n_train], all_rois[:n_train], all_labels[:n_train],
    all_diffs[:n_train], all_bboxes[:n_train])

test_set = rcnn.Dataset(
    img_paths[n_train:], all_rois[n_train:], all_labels[n_train:],
    all_diffs[n_train:], all_bboxes[n_train:])
```

In [24]:

```
## check the results
image, crops, roi_bboxes, labels, diffs, bboxes, img_path = train_set[0]

plt.rcParams['figure.figsize'] = [5, 3]
fig, ax = plt.subplots()
ax.imshow(crops[3])
plt.show()

print(labels[3])
```



1

In [27]:

```
train_loader = DataLoader(
    train_set, batch_size=2, collate_fn=train_set.collate_fn, drop_last=True)
```

In [28]:

```
test_loader = DataLoader(
    test_set, batch_size=2, collate_fn=test_set.collate_fn, drop_last=True)
```

Define the backbone

In [29]:

```
backbone = utils.get_backbone()
```

```
Downloading: "https://download.pytorch.org/models/vgg16-397923af.pth" to /root/.cache/torch/hub/checkpoints/vgg16-397923af.pth
```

In [30]:

```
device = 'cuda' if torch.cuda.is_available() else 'cpu'  
device
```

Out[30]:

```
'cuda'
```

In [31]:

```
backbone.eval().to(device)
```

Out[31]:

```
VGG(  
    (features): Sequential(  
        (0): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (1): ReLU(inplace=True)  
        (2): Conv2d(64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (3): ReLU(inplace=True)  
        (4): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (5): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (6): ReLU(inplace=True)  
        (7): Conv2d(128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (8): ReLU(inplace=True)  
        (9): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (10): Conv2d(128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (11): ReLU(inplace=True)  
        (12): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (13): ReLU(inplace=True)  
        (14): Conv2d(256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (15): ReLU(inplace=True)  
        (16): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (17): Conv2d(256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (18): ReLU(inplace=True)  
        (19): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (20): ReLU(inplace=True)  
        (21): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (22): ReLU(inplace=True)  
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
        (24): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (25): ReLU(inplace=True)  
        (26): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (27): ReLU(inplace=True)  
        (28): Conv2d(512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1))  
        (29): ReLU(inplace=True)  
        (30): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)  
    )  
    (avgpool): AdaptiveAvgPool2d(output_size=(7, 7))  
    (classifier): Sequential()  
)
```

Train the model

In [32]:

```
model = rcnn.BaseModel(backbone).to(device)  
criterion = model.calc_loss  
optimizer = optim.SGD(model.parameters(), lr=1e-3)
```

In [33]:

```
n_epochs = 5

## log for plot training and validation metrics
log = Report(n_epochs)
tstart = time.time()
for epoch in range(n_epochs):

    _n = len(train_loader)
    for i, inputs in enumerate(train_loader):
        loss, loc_loss, regr_loss, accs = rcnn.train_batch(
            inputs, model, optimizer, criterion)
        pos = (epoch + (i + 1)/_n)
        log.record(pos,
                   trn_loss=loss.item(),
                   trn_loc_loss=loc_loss,
                   trn_regr_loss=regr_loss,
                   trn_acc=accs.mean(),
                   end='\r')

    _n = len(test_loader)
    for i, inputs in enumerate(test_loader):
        _, _, _, _, loc_loss, regr_loss, accs = rcnn.validate_batch(
            inputs, model, criterion)
        pos = (epoch + (i + 1)/_n)
        log.record(
            pos,
            val_loss=loss.item(),
            val_loc_loss=loc_loss,
            val_regr_loss=regr_loss,
            val_acc=accs.mean(),
            end='\r')

tend = time.time()

EPOCH: 5.000      val_loss: 0.032 val_loc_loss: 0.032      val_regr_loss: 0.000      val_acc:
```

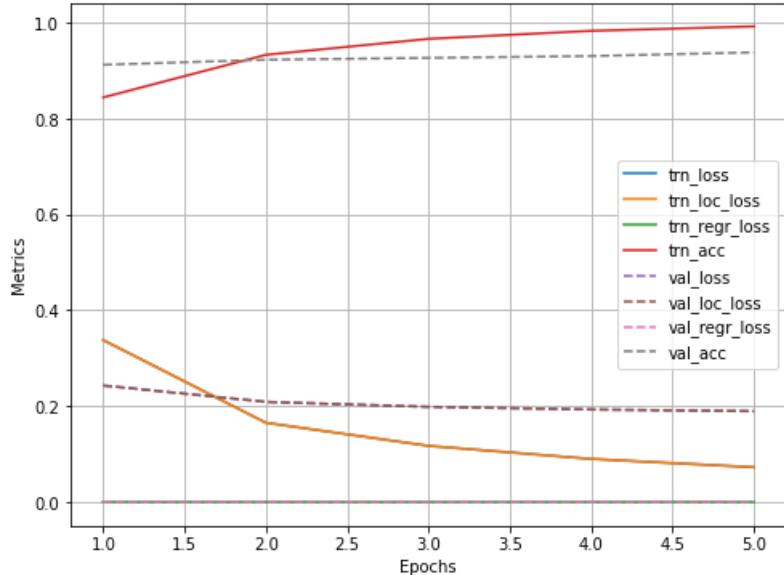
1.000 (219.84s - 0.00s remaining))

In [34]:

```
print('Time elapsed = %.2f min' % ((tend - tstart)/60))
log.plot_epochs()
```

100%|██████████| 106/106 [00:00<00:00, 1784.15it/s]

Time elapsed = 3.66 min



In [35]:

```
def show_results(i):
    img_name = all_img_names[i]
    img_path = imgs_path + img_name
    img = Image.open(img_path)
    np_image = np.array(img)

    ## reverse the code from before
    regions = utils.get_simple_regions()
    input = []
    crops = []
    for region in regions:
        x, y, w, h = region
        x0, y0, x1, y1 = x, y, x + w, y + h
        crop = np_image[y0:y1, x0:x1]
        crops.append(crop)

    newsize = (224, 224)
    crops = [Image.fromarray(crop, 'RGB') for crop in crops]
    crops = [crop.resize(newsize) for crop in crops]
    crops = [utils.preprocess_image(crop)[None] for crop in crops]

    input = torch.cat(crops).to(device)
    with torch.no_grad():
        model.eval()
        probs, diffs = model(input)
        probs = torch.nn.functional.softmax(probs, -1)
        confs, classes = torch.max(probs, -1)

    regions = np.array(regions)
    confs, classes, probs, diffs = [
        tensor.detach().cpu().numpy() for tensor in [confs, classes, probs, diffs]]

    # TODO: use nms to lower the recall
    ## adding predicted diffs
    detected_bboxes = (regions + diffs).astype(np.uint16)

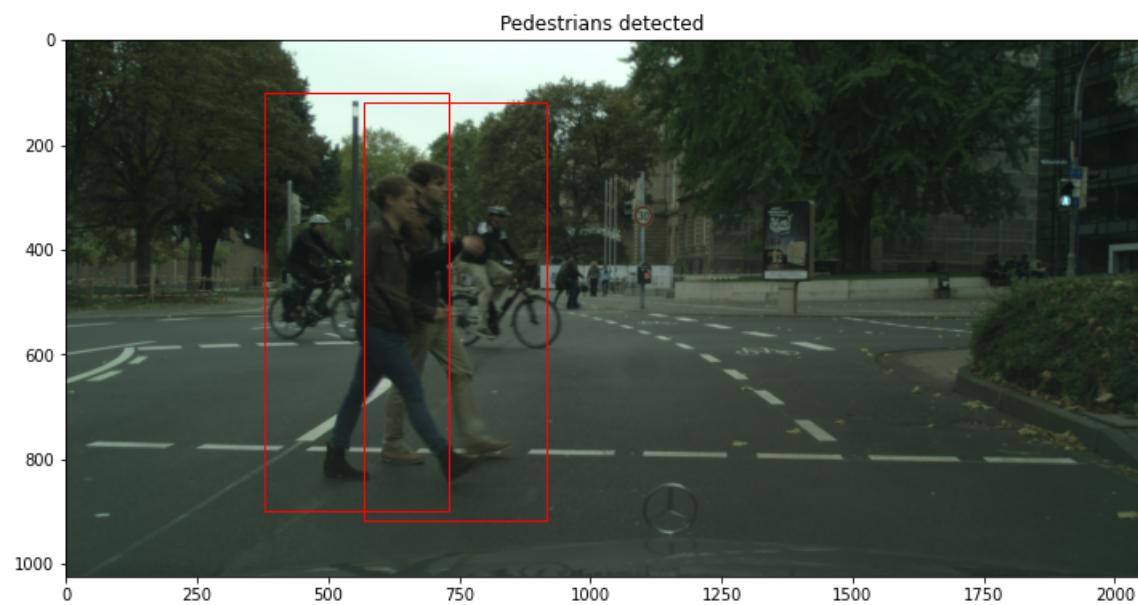
    plt.rcParams['figure.figsize'] = [12, 8]
    fig, ax = plt.subplots()

    plt.title('Pedestrians detected')
    ax.imshow(img)

    classes = classes.tolist()
    for i, bbox in enumerate(detected_bboxes):
        if classes[i] == 1:
            rect = patches.Rectangle(
                (bbox[0], bbox[1]), bbox[2], bbox[3],
                linewidth=1, edgecolor='r', facecolor='none')
            ax.add_patch(rect)
    plt.show()
```

In [36]:

```
show_results(0)
```



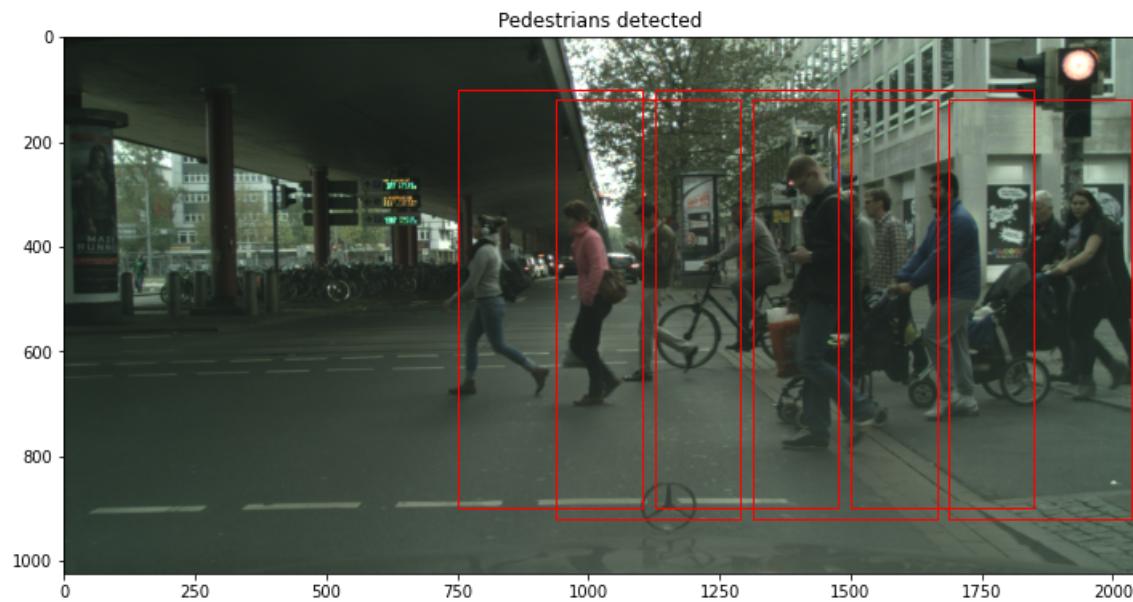
In [37]:

```
show_results(1)
```



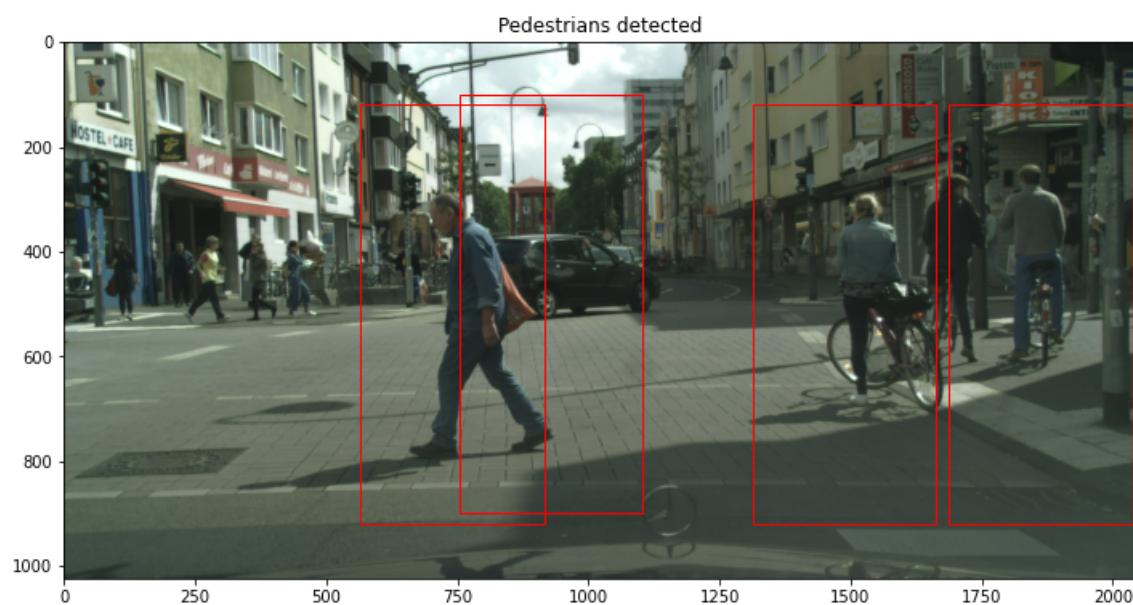
In [38]:

```
show_results(2)
```



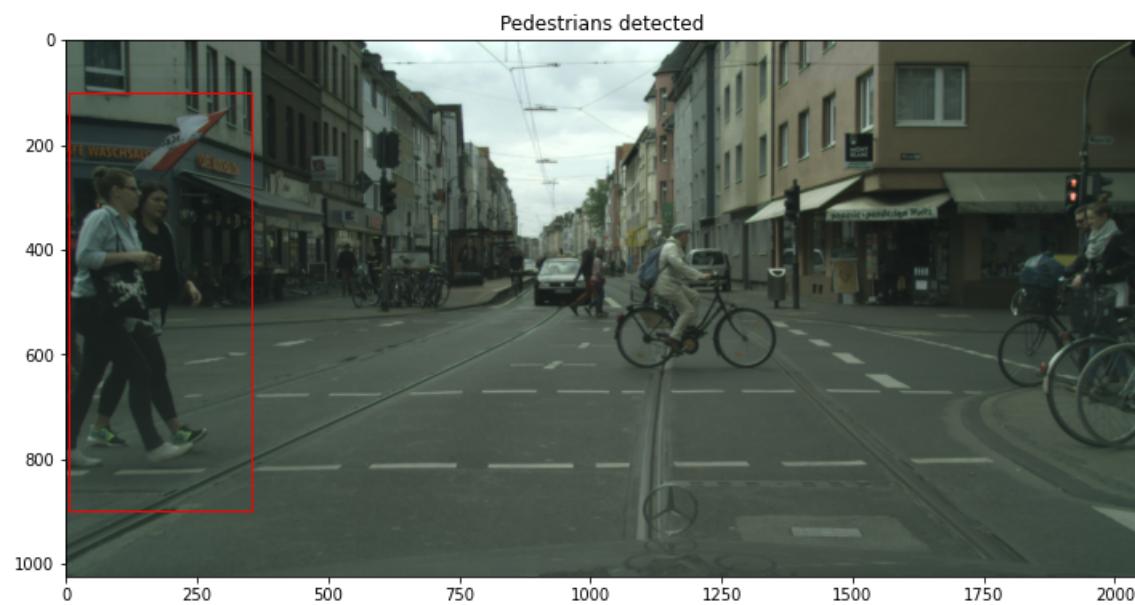
In [39]:

```
show_results(3)
```



In [53]:

```
show_results(4)
```



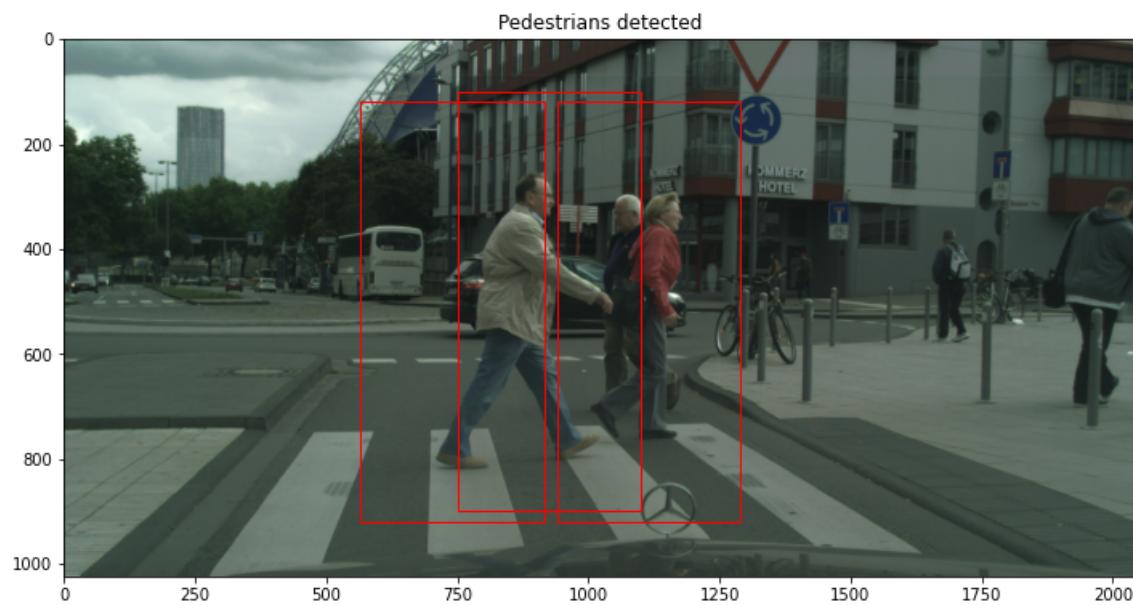
In [54]:

```
show_results(5)
```



In [40]:

```
show_results(6)
```



In [41]:

```
show_results(7)
```



In [42]:

```
show_results(8)
```



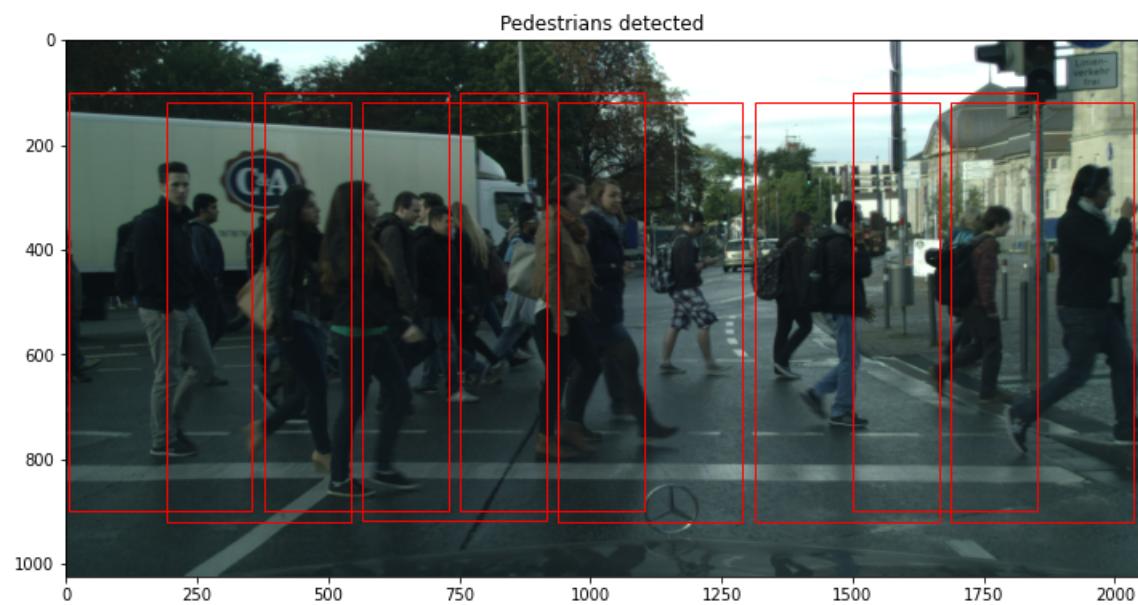
In [43]:

```
show_results(9)
```



In [44]:

```
show_results(10)
```



In []: